

Assignment 6 :

Problem Statement

1. Execution at terminal to create the database and table and import the table in the HDFS at location Sqoopout folder :

```
mysql> create database simplidb;  
Query OK, 1 row affected (0.11 sec)
```

```
mysql> show databases;
```

```
+-----+  
| Database      |  
+-----+  
| hive_mysql_db |  
| information_schema |  
| mysql         |  
| performance_schema |  
| simplidb      |  
| sqoopbatch    |  
| sys           |  
+-----+
```

7 rows in set (0.00 sec)

```
mysql> use simplidb;
```

Database changed

```
mysql> Create table Person(  
-> person_id INT NOT NULL,  
-> lname varchar(20),  
-> fname varchar(20),  
-> area varchar(20),  
-> city varchar(20),  
-> PRIMARY KEY(person_id)  
-> );
```

Query OK, 0 rows affected (0.09 sec)

```
mysql> insert into Person(person_id,lname,fname,area,city) values  
(1,"Abhishek","Ram","Patna","Bihar");
```

Query OK, 1 row affected (0.06 sec)

```
mysql> insert into Person(person_id,lname,fname,area,city) values
(2,"Tanya","Priya","Whitefield","Banglore");
Query OK, 1 row affected (0.08 sec)
```

```
mysql> insert into Person(person_id,lname,fname,area,city) values (3,"James",
"Brown","New York","United States");
Query OK, 1 row affected (0.09 sec)
```

```
mysql> insert into Person(person_id,lname,fname,area,city) values (4,"John",
"Miller","Los Angeles","United States");
Query OK, 1 row affected (0.04 sec)
```

```
mysql> select * from Person;
```

person_id	lname	fname	area	city
1	Abhishek	Ram	Patna	Bihar
2	Tanya	Priya	Whitefield	Banglore
3	James	Brown	New York	United States
4	John	Miller	Los Angeles	United States

4 rows in set (0.00 sec)

```
mysql> Update Person set lname="Shyam" where person_id =1;
Query OK, 1 row affected (0.09 sec)
Rows matched: 1  Changed: 1  Warnings: 0
```

```
mysql> select * from Person;
```

person_id	lname	fname	area	city
1	Shyam	Ram	Patna	Bihar
2	Tanya	Priya	Whitefield	Banglore
3	James	Brown	New York	United States
4	John	Miller	Los Angeles	United States

4 rows in set (0.00 sec)

```
mysql> ^C
mysql> exit
Bye
```

```
[acadgild@localhost ~]$ sqoop import --connect  
jdbc:mysql://localhost:3306/simplidb --table Person --target-dir /sqoopout -m 1  
--username root -P
```

Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/../hcatalog does not exist! HCatalog jobs will fail.

Please set \$HCAT_HOME to the root of your HCatalog installation.

Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/../accumulo does not exist! Accumulo imports will fail.

Please set \$ACCUMULO_HOME to the root of your Accumulo installation.

18/07/04 02:01:16 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7

Enter password:

18/07/04 02:01:21 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.

18/07/04 02:01:21 INFO tool.CodeGenTool: Beginning code generation

Wed Jul 04 02:01:22 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/04 02:01:23 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1

18/07/04 02:01:23 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1

18/07/04 02:01:23 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5

Note: /tmp/sqoop-acadgild/compile/62bcc1b5c8d16182e758da71bf74858b/Person.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

18/07/04 02:01:32 INFO orm.CompilationManager: Writing jar file:

/tmp/sqoop-acadgild/compile/62bcc1b5c8d16182e758da71bf74858b/Person.jar

18/07/04 02:01:32 WARN manager.MySQLManager: It looks like you are importing from mysql.

18/07/04 02:01:32 WARN manager.MySQLManager: This transfer can be faster! Use the --direct

18/07/04 02:01:32 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.

18/07/04 02:01:32 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)

18/07/04 02:01:32 INFO mapreduce.ImportJobBase: Beginning import of Person

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hbase/hbase-1.4.4/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

18/07/04 02:01:32 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/07/04 02:01:32 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

18/07/04 02:01:33 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

18/07/04 02:01:34 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

Wed Jul 04 02:01:40 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/04 02:01:40 INFO db.DBInputFormat: Using read committed transaction isolation

18/07/04 02:01:41 INFO mapreduce.JobSubmitter: number of splits:1

18/07/04 02:01:41 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1530640014524_0001

18/07/04 02:01:42 INFO impl.YarnClientImpl: Submitted application application_1530640014524_0001

18/07/04 02:01:43 INFO mapreduce.Job: The url to track the job:

http://localhost:8088/proxy/application_1530640014524_0001/

18/07/04 02:01:43 INFO mapreduce.Job: Running job: job_1530640014524_0001

18/07/04 02:02:04 INFO mapreduce.Job: Job job_1530640014524_0001 running in uber mode : false

18/07/04 02:02:04 INFO mapreduce.Job: map 0% reduce 0%

18/07/04 02:02:15 INFO mapreduce.Job: map 100% reduce 0%

18/07/04 02:02:16 INFO mapreduce.Job: Job job_1530640014524_0001 completed successfully

18/07/04 02:02:16 INFO mapreduce.Job: Counters: 30

File System Counters

FILE: Number of bytes read=0

FILE: Number of bytes written=128195
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=87
HDFS: Number of bytes written=135
HDFS: Number of read operations=4
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Other local map tasks=1
Total time spent by all maps in occupied slots (ms)=7241
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=7241
Total vcore-milliseconds taken by all map tasks=7241
Total megabyte-milliseconds taken by all map tasks=7414784

Map-Reduce Framework

Map input records=4
Map output records=4
Input split bytes=87
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=146
CPU time spent (ms)=1680
Physical memory (bytes) snapshot=75149312
Virtual memory (bytes) snapshot=347095040
Total committed heap usage (bytes)=21573632

File Input Format Counters

Bytes Read=0

File Output Format Counters

Bytes Written=135

18/07/04 02:02:16 INFO mapreduce.ImportJobBase: Transferred 135 bytes in 42.4646 seconds (3.1791 bytes/sec)

18/07/04 02:02:16 INFO mapreduce.ImportJobBase: Retrieved 4 records.

[acadgild@localhost ~]\$ hdfs dfs -ls /

18/07/04 02:02:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 4 items

drwxr-xr-x	- acadgild supergroup	0 2018-05-16 13:49 /hbase
drwxr-xr-x	- acadgild supergroup	0 2018-07-04 02:02 /sqoopout

```

drwx-wx-wx - acadgild supergroup      0 2018-05-16 13:42 /tmp
drwxr-xr-x - acadgild supergroup      0 2018-06-24 01:28 /user
[acadgild@localhost ~]$ hdfs dfs -ls /user/
18/07/04 02:03:03 WARN util.NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x - acadgild supergroup      0 2018-06-24 01:29 /user/acadgild
[acadgild@localhost ~]$ hdfs dfs -ls /user/acadgild
18/07/04 02:03:13 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x - acadgild supergroup      0 2018-06-29 20:47
/user/acadgild/hadoopdata
[acadgild@localhost ~]$ hdfs dfs -ls /sqoopout
18/07/04 02:03:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 acadgild supergroup      0 2018-07-04 02:02 /sqoopout/_SUCCESS
-rw-r--r--  1 acadgild supergroup    135 2018-07-04 02:02
/sqoopout/part-m-00000
[acadgild@localhost ~]$ hdfs dfs -cat /sqoopout/part-m-00000
18/07/04 02:04:25 WARN util.NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Patna,Bihar,Ram,Shyam,1
Whitefield,Banglore,Priya,Tanya,2
New York,United States,Brown,James,3
Los Angeles,United States,Miller,John,4

```

Note: Now by now the table data has been imported from mysql relational database to the HDFS at location /Sqoopout folder.

Task 1

Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table

Solution 1:

Here, we will create a new database with name : **sqoop_export1**
and a table : Person1 with same column_names of table as Person table in Simplidb database;

Therefore, we are exported from HDFS to Relational Table Person using sqoop.

Terminal Execution :

```
[acadgild@localhost ~]$ mysql -u root -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 13
Server version: 8.0.11 MySQL Community Server - GPL
```

Copyright (c) 2000, 2018, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> Show databases;
+-----+
| Database          |
+-----+
| hive_mysql_db     |
| information_schema|
| mysql             |
| performance_schema|
| simplidb          |
| sqoopbatch        |
| sys               |
+-----+
7 rows in set (0.00 sec)
```

```
mysql> create database sqoop_export1;
Query OK, 1 row affected (0.10 sec)
```

```
mysql> show databases;
```

```
+-----+
| Database      |
+-----+
| hive_mysql_db |
| information_schema |
| mysql         |
| performance_schema |
| simplidb      |
| sqoop_export1 |
| sqoopbatch    |
| sys          |
+-----+
```

```
8 rows in set (0.01 sec)
```

```
mysql> show tables in sqoop_export1;
Empty set (0.00 sec)
```

```
mysql> use sqoop_export1;
Database changed
```

```
mysql> create table Person1(
-> lname varchar(20),
-> fname varchar(20),
-> person_id INT NOT NULL,
-> area varchar(20),
-> city varchar(20),
-> PRIMARY KEY(person_id)
-> );
```

```
Query OK, 0 rows affected (0.14 sec)
```

```
mysql> show tables in sqoop_export1;
```

```
+-----+
| Tables_in_sqoop_export1 |
+-----+
| Person1                 |
+-----+
```

```
1 row in set (0.01 sec)
```

```
mysql> select * from Person1;
Empty set (0.01 sec)
```

```
mysql> exit
```

```
Bye
```

```
[acadgild@localhost ~]$ hdfs dfs -ls /sqoopout
```


18/07/04 02:33:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

```
-rw-r--r--  1 acadgild supergroup      0 2018-07-04 02:02 /sqoopout/_SUCCESS
-rw-r--r--  1 acadgild supergroup 135 2018-07-04 02:02 /sqoopout/part-m-00000
```

[acadgild@localhost ~]\$ hdfs dfs -cat /sqoopout/part-m-00000

18/07/04 02:33:32 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Patna,Bihar,Ram,Shyam,1

Whitefield,Bangalore,Priya,Tanya,2

New York,United States,Brown,James,3

Los Angeles,United States,Miller,John,4

[acadgild@localhost ~]\$ sqoop export --connect jdbc:mysql://localhost:3306/sqoop_export1 --table Person1 -m 1 --export-dir /sqoopout/ --username root -P

Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/./hcatalog does not exist!
HCatalog jobs will fail.

Please set \$HCAT_HOME to the root of your HCatalog installation.

Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/./accumulo does not exist!
Accumulo imports will fail.

Please set \$ACCUMULO_HOME to the root of your Accumulo installation.

18/07/04 02:49:30 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7

Enter password:

18/07/04 02:49:38 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.

18/07/04 02:49:38 INFO tool.CodeGenTool: Beginning code generation

Wed Jul 04 02:49:39 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/04 02:49:40 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person1` AS t LIMIT 1

18/07/04 02:49:40 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person1` AS t LIMIT 1

18/07/04 02:49:40 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5

Note: /tmp/sqoop-acadgild/compile/b7a7f7c45e106418dcec73bdf5c096cd/Person1.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

18/07/04 02:49:44 INFO orm.CompilationManager: Writing jar file:

/tmp/sqoop-acadgild/compile/b7a7f7c45e106418dcec73bdf5c096cd/Person1.jar

18/07/04 02:49:44 INFO mapreduce.ExportJobBase: Beginning export of Person1

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hbase/hbase-1.4.4/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLogg

erBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

18/07/04 02:49:44 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/07/04 02:49:44 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

18/07/04 02:49:46 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative

18/07/04 02:49:46 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

18/07/04 02:49:46 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

18/07/04 02:49:47 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

18/07/04 02:49:51 INFO input.FileInputFormat: Total input paths to process : 1

18/07/04 02:49:51 INFO input.FileInputFormat: Total input paths to process : 1

18/07/04 02:49:51 INFO mapreduce.JobSubmitter: number of splits:1

18/07/04 02:49:51 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

18/07/04 02:49:52 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1530640014524_0003

18/07/04 02:49:52 INFO impl.YarnClientImpl: Submitted application
application_1530640014524_0003

18/07/04 02:49:53 INFO mapreduce.Job: The url to track the job:

http://localhost:8088/proxy/application_1530640014524_0003/

18/07/04 02:49:53 INFO mapreduce.Job: Running job: job_1530640014524_0003

18/07/04 02:50:10 INFO mapreduce.Job: Job job_1530640014524_0003 running in uber mode : false

18/07/04 02:50:10 INFO mapreduce.Job: map 0% reduce 0%

18/07/04 02:50:20 INFO mapreduce.Job: map 100% reduce 0%

18/07/04 02:50:21 INFO mapreduce.Job: Job job_1530640014524_0003 completed successfully

18/07/04 02:50:22 INFO mapreduce.Job: Counters: 30

File System Counters

FILE: Number of bytes read=0

FILE: Number of bytes written=127884

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=262

HDFS: Number of bytes written=0

HDFS: Number of read operations=4

HDFS: Number of large read operations=0

HDFS: Number of write operations=0

Job Counters

Launched map tasks=1

Data-local map tasks=1

Total time spent by all maps in occupied slots (ms)=7540

Total time spent by all reduces in occupied slots (ms)=0

Total time spent by all map tasks (ms)=7540

Total vcore-milliseconds taken by all map tasks=7540

Total megabyte-milliseconds taken by all map tasks=7720960
Map-Reduce Framework
Map input records=4
Map output records=4
Input split bytes=124
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=139
CPU time spent (ms)=1480
Physical memory (bytes) snapshot=73748480
Virtual memory (bytes) snapshot=347480064
Total committed heap usage (bytes)=17199104
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=0

18/07/04 02:50:22 INFO mapreduce.ExportJobBase: Transferred 262 bytes in 35.1718 seconds (7.4491 bytes/sec)

18/07/04 02:50:22 INFO mapreduce.ExportJobBase: Exported 4 records.

[acadgild@localhost ~]\$ mysql -u root -p

Enter password:

Welcome to the MySQL monitor. Commands end with ; or \g.

Your MySQL connection id is 17

Server version: 8.0.11 MySQL Community Server - GPL

Copyright (c) 2000, 2018, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> show databases;

```
+-----+
| Database          |
+-----+
| hive_mysql_db     |
| information_schema|
| mysql             |
| performance_schema|
| simplidb          |
| sqoop_export1     |
| sqoopbatch        |
| sys               |
+-----+
```

8 rows in set (0.00 sec)

```
mysql> use sqoop_export1;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

Database changed

```
mysql> select * from Person1;
```

```
+-----+-----+-----+-----+-----+
| lname | fname | person_id | area      | city      |
+-----+-----+-----+-----+-----+
| Shyam | Ram   | 1         | Patna     | Bihar     |
| Tanya | Priya | 2         | Whitefield | Bangalore  |
| James | Brown | 3         | New York  | United States |
| John  | Miller | 4         | Los Angeles | United States |
+-----+-----+-----+-----+-----+
4 rows in set (0.00 sec)
```

```
mysql>
```

Task2

Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table with parameter person_id =3.

Solution 2 :

Here, we will use the database with name : **sqoop_export1**

and a table : Person2 with same column_names of table as Person table in Simplidb database;

Therefore, we are exported from HDFS to Relational Table Person using sqoop that has person_id = 3.

Terminal Execution :

```
[acadgild@localhost ~]$ mysql -u root -p
```

```
Enter password:
```

```
Welcome to the MySQL monitor.  Commands end with ; or \g.
```

```
Your MySQL connection id is 16
```

```
Server version: 8.0.11 MySQL Community Server - GPL
```

Copyright (c) 2000, 2018, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its

affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> use sqoop_export1;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

Database changed

```
mysql> show tables;
```

```
+-----+
| Tables_in_sqoop_export1 |
+-----+
| Person1                  |
| Person2                  |
| Person3                  |
+-----+
```

3 rows in set (0.01 sec)

```
mysql> create table PersonExport(
```

```
-> person_id INT NOT NULL,
-> lname varchar(20),
-> fname varchar(20),
-> area varchar(20),
-> city varchar(20),
-> Primary Key (person_id));
```

Query OK, 0 rows affected (0.12 sec)

```
mysql> show tables;
```

```
+-----+
| Tables_in_sqoop_export1 |
+-----+
| Person1                  |
| Person2                  |
| Person3                  |
| PersonExport              |
+-----+
```

4 rows in set (0.00 sec)

```
mysql> select * from PersonExport;
```

Empty set (0.01 sec)

```
mysql> exit
```

Bye

```
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost:3306/simplidb --table Person
--target-dir /SQOOPOUT/ --split-by person_id --username root -P
Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/./hcatalog does not exist!
```

HCatalog jobs will fail.

Please set \$HCAT_HOME to the root of your HCatalog installation.

Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/./accumulo does not exist!
Accumulo imports will fail.

Please set \$ACCUMULO_HOME to the root of your Accumulo installation.

18/07/04 22:38:56 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7

Enter password:

18/07/04 22:39:01 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.

18/07/04 22:39:01 INFO tool.CodeGenTool: Beginning code generation

Wed Jul 04 22:39:01 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/04 22:39:02 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1

18/07/04 22:39:02 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `Person` AS t LIMIT 1

18/07/04 22:39:02 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5

Note: /tmp/sqoop-acadgild/compile/168a490dec92fa7825b0f4fff31c1853/Person.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

18/07/04 22:39:06 INFO orm.CompilationManager: Writing jar file:

/tmp/sqoop-acadgild/compile/168a490dec92fa7825b0f4fff31c1853/Person.jar

18/07/04 22:39:06 WARN manager.MySQLManager: It looks like you are importing from mysql.

18/07/04 22:39:06 WARN manager.MySQLManager: This transfer can be faster! Use the --direct

18/07/04 22:39:06 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.

18/07/04 22:39:06 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)

18/07/04 22:39:06 INFO mapreduce.ImportJobBase: Beginning import of Person

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hbase/hbase-1.4.4/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

18/07/04 22:39:06 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/07/04 22:39:06 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

18/07/04 22:39:07 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

18/07/04 22:39:07 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

Wed Jul 04 22:39:12 IST 2018 WARN: Establishing SSL connection without server's identity

verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/04 22:39:12 INFO db.DBInputFormat: Using read committed transaction isolation

18/07/04 22:39:12 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`person_id`), MAX(`person_id`) FROM `Person`

18/07/04 22:39:12 INFO db.IntegerSplitter: Split size: 0; Num splits: 4 from: 1 to: 4

18/07/04 22:39:12 INFO mapreduce.JobSubmitter: number of splits:4

18/07/04 22:39:12 INFO mapreduce.JobSubmitter: Submitting tokens for job:

job_1530722670909_0002

18/07/04 22:39:13 INFO impl.YarnClientImpl: Submitted application

application_1530722670909_0002

18/07/04 22:39:13 INFO mapreduce.Job: The url to track the job:

http://localhost:8088/proxy/application_1530722670909_0002/

18/07/04 22:39:13 INFO mapreduce.Job: Running job: job_1530722670909_0002

18/07/04 22:39:27 INFO mapreduce.Job: Job job_1530722670909_0002 running in uber mode : false

18/07/04 22:39:27 INFO mapreduce.Job: map 0% reduce 0%

18/07/04 22:39:58 INFO mapreduce.Job: map 25% reduce 0%

18/07/04 22:39:59 INFO mapreduce.Job: map 50% reduce 0%

18/07/04 22:40:00 INFO mapreduce.Job: map 75% reduce 0%

18/07/04 22:40:01 INFO mapreduce.Job: map 100% reduce 0%

18/07/04 22:40:01 INFO mapreduce.Job: Job job_1530722670909_0002 completed successfully

18/07/04 22:40:01 INFO mapreduce.Job: Counters: 31

File System Counters

FILE: Number of bytes read=0

FILE: Number of bytes written=512780

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=449

HDFS: Number of bytes written=135

HDFS: Number of read operations=16

HDFS: Number of large read operations=0

HDFS: Number of write operations=8

Job Counters

Killed map tasks=1

Launched map tasks=4

Other local map tasks=4

Total time spent by all maps in occupied slots (ms)=112587

Total time spent by all reduces in occupied slots (ms)=0

Total time spent by all map tasks (ms)=112587

Total vcore-milliseconds taken by all map tasks=112587

Total megabyte-milliseconds taken by all map tasks=115289088

Map-Reduce Framework

Map input records=4

Map output records=4

Input split bytes=449

Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=1692
CPU time spent (ms)=7390
Physical memory (bytes) snapshot=302428160
Virtual memory (bytes) snapshot=1391083520
Total committed heap usage (bytes)=86020096
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=135

18/07/04 22:40:01 INFO mapreduce.ImportJobBase: Transferred 135 bytes in 53.8703 seconds (2.506 bytes/sec)

18/07/04 22:40:01 INFO mapreduce.ImportJobBase: Retrieved 4 records.

[acadgild@localhost ~]\$ hdfs dfs -ls /SQOOPOUT/

18/07/04 22:43:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 5 items

-rw-r--r--	1	acadgild supergroup	0	2018-07-04 22:39	/SQOOPOUT/_SUCCESS
-rw-r--r--	1	acadgild supergroup	24	2018-07-04 22:39	/SQOOPOUT/part-m-00000
-rw-r--r--	1	acadgild supergroup	34	2018-07-04 22:39	/SQOOPOUT/part-m-00001
-rw-r--r--	1	acadgild supergroup	37	2018-07-04 22:39	/SQOOPOUT/part-m-00002
-rw-r--r--	1	acadgild supergroup	40	2018-07-04 22:39	/SQOOPOUT/part-m-00003

[acadgild@localhost ~]\$ hdfs dfs -cat /SQOOPOUT/part-m-00002

18/07/04 22:44:12 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

New York,United States,Brown,James,3

[acadgild@localhost ~]\$ sqoop export --connect jdbc:mysql://localhost:3306/sqoop_export1 --table

PersonExport --export-dir /SQOOPOUT/part-m-00002 --username root -P

Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/./hcatalog does not exist!
HCatalog jobs will fail.

Please set \$HCAT_HOME to the root of your HCatalog installation.

Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/./accumulo does not exist!
Accumulo imports will fail.

Please set \$ACCUMULO_HOME to the root of your Accumulo installation.

18/07/04 22:47:25 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7

Enter password:

18/07/04 22:47:31 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.

18/07/04 22:47:31 INFO tool.CodeGenTool: Beginning code generation

Wed Jul 04 22:47:31 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to

explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/04 22:47:32 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `PersonExport` AS t LIMIT 1

18/07/04 22:47:32 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `PersonExport` AS t LIMIT 1

18/07/04 22:47:32 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5

Note: /tmp/sqoop-acadgild/compile/9b67cb31fd26ee856e85ec798f15db9c/PersonExport.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

18/07/04 22:47:35 INFO orm.CompilationManager: Writing jar file:

/tmp/sqoop-acadgild/compile/9b67cb31fd26ee856e85ec798f15db9c/PersonExport.jar

18/07/04 22:47:35 INFO mapreduce.ExportJobBase: Beginning export of PersonExport

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hbase/hbase-1.4.4/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

18/07/04 22:47:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/07/04 22:47:36 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

18/07/04 22:47:38 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative

18/07/04 22:47:38 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

18/07/04 22:47:38 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

18/07/04 22:47:38 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

18/07/04 22:47:41 INFO input.FileInputFormat: Total input paths to process : 1

18/07/04 22:47:41 INFO input.FileInputFormat: Total input paths to process : 1

18/07/04 22:47:41 INFO mapreduce.JobSubmitter: number of splits:4

18/07/04 22:47:41 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

18/07/04 22:47:42 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1530722670909_0003

18/07/04 22:47:42 INFO impl.YarnClientImpl: Submitted application application_1530722670909_0003

18/07/04 22:47:42 INFO mapreduce.Job: The url to track the job:

http://localhost:8088/proxy/application_1530722670909_0003/

18/07/04 22:47:42 INFO mapreduce.Job: Running job: job_1530722670909_0003

18/07/04 22:47:55 INFO mapreduce.Job: Job job_1530722670909_0003 running in uber mode : false

18/07/04 22:47:55 INFO mapreduce.Job: map 0% reduce 0%

18/07/04 22:48:24 INFO mapreduce.Job: map 50% reduce 0%

18/07/04 22:48:25 INFO mapreduce.Job: map 100% reduce 0%
18/07/04 22:48:25 INFO mapreduce.Job: Job job_1530722670909_0003 completed successfully
18/07/04 22:48:25 INFO mapreduce.Job: Counters: 30

File System Counters

FILE: Number of bytes read=0
FILE: Number of bytes written=511648
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=670
HDFS: Number of bytes written=0
HDFS: Number of read operations=19
HDFS: Number of large read operations=0
HDFS: Number of write operations=0

Job Counters

Launched map tasks=4
Data-local map tasks=4
Total time spent by all maps in occupied slots (ms)=110284
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=110284
Total vcore-milliseconds taken by all map tasks=110284
Total megabyte-milliseconds taken by all map tasks=112930816

Map-Reduce Framework

Map input records=1
Map output records=1
Input split bytes=556
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=1524
CPU time spent (ms)=6320
Physical memory (bytes) snapshot=295948288
Virtual memory (bytes) snapshot=1389920256
Total committed heap usage (bytes)=68976640

File Input Format Counters

Bytes Read=0

File Output Format Counters

Bytes Written=0

18/07/04 22:48:25 INFO mapreduce.ExportJobBase: Transferred 670 bytes in 47.931 seconds (13.9784 bytes/sec)

18/07/04 22:48:25 INFO mapreduce.ExportJobBase: Exported 1 records.

[acadgild@localhost ~]\$ mysql -u root -p

Enter password:

Welcome to the MySQL monitor. Commands end with ; or \g.

Your MySQL connection id is 28

Server version: 8.0.11 MySQL Community Server - GPL

Copyright (c) 2000, 2018, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> use sqoop_export1;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

```
Database changed
mysql> show tables;
+-----+
| Tables_in_sqoop_export1 |
+-----+
| Person1                  |
| Person2                  |
| Person3                  |
| PersonExport              |
+-----+
4 rows in set (0.00 sec)
```

```
mysql> select * from PersonExport;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area   | city      |
+-----+-----+-----+-----+-----+
|          3 | James | Brown | New York | United States |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

Explanation:

Since, we while importing from mysql to hdfs directory /SQOOPOUT/ we have used below command :

```
sqoop import --connect jdbc:mysql://localhost:3306/simplidb --table Person --target-dir /SQOOPOUT/ --split-by person_id --username root -P
```

here, we have used split-by person_id without defining the mappers. Hence, each person_id gets stored in their respective mapper. And according to our execution, record with person_id = 3 lies in the '/SQOOPOUT/part-m-00002' path in HDFS.

So, we use below statement to directly export content of '/SQOOPOUT/part-m-00002' path in HDFS which contains the record containing person_id=3 to the mysql table PersonExport :

```
[acadgild@localhost ~]$ sqoop export --connect jdbc:mysql://localhost:3306/sqoop_export1
```

--table PersonExport --export-dir /SQOOPOUT/part-m-00002 --username root -P

Hence, we are successfully able to export detail of person with id =3 from HDFS to MySql relational database sqoop_export1 and table PersonExport.

2nd Method of doing task 2:

That is exporting the details of record person_id=3 in HDFS to MySQL table PersonExportNew in sqoop_export1 database.

Terminal Execution:

```
mysql> create table PersonExportNew( person_id INT NOT NULL, lname varchar(20), fname
varchar(20), area varchar(20), city varchar(20), Primary Key (person_id));
Query OK, 0 rows affected (0.13 sec)
```

```
mysql> select * from PersonExportNew;
set (0.00 sec) Empty
```

```
mysql> show table from sqoop_export1;
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to
your MySQL server version for the right syntax to use near 'from sqoop_export1' at line 1
mysql> show tables;
```

```
+-----+
| Tables_in_sqoop_export1 |
+-----+
| Person1                |
| Person2                |
| Person3                |
| PersonExport           |
| PersonExportNew        |
+-----+
5 rows in set (0.00 sec)
```

```
mysql> exit
Bye
```

```
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost:3306/simplidb --target-dir
/SQOOPOUT1/ -m 1 --query 'select * from Person where person_id=3 and $CONDITIONS'
--username root -P
Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/./hcatalog does not exist!
HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
```

Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0/./accumulo does not exist!
Accumulo imports will fail.

Please set \$ACCUMULO_HOME to the root of your Accumulo installation.

18/07/04 23:12:43 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7

Enter password:

18/07/04 23:12:48 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.

18/07/04 23:12:48 INFO tool.CodeGenTool: Beginning code generation

Wed Jul 04 23:12:48 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/04 23:12:49 INFO manager.SqlManager: Executing SQL statement: select * from Person where person_id=3 and (1 = 0)

18/07/04 23:12:49 INFO manager.SqlManager: Executing SQL statement: select * from Person where person_id=3 and (1 = 0)

18/07/04 23:12:49 INFO manager.SqlManager: Executing SQL statement: select * from Person where person_id=3 and (1 = 0)

18/07/04 23:12:49 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is
/home/acadgild/install/hadoop/hadoop-2.6.5

Note: /tmp/sqoop-acadgild/compile/70bd85f878167801b95972008b7561ec/QueryResult.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

18/07/04 23:12:52 INFO orm.CompilationManager: Writing jar file:

/tmp/sqoop-acadgild/compile/70bd85f878167801b95972008b7561ec/QueryResult.jar

18/07/04 23:12:53 INFO mapreduce.ImportJobBase: Beginning query import.

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in

[jar:file:/home/acadgild/install/hbase/hbase-1.4.4/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

18/07/04 23:12:53 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/07/04 23:12:53 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use
mapreduce.job.jar

18/07/04 23:12:54 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use
mapreduce.job.maps

18/07/04 23:12:54 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

Wed Jul 04 23:12:58 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/04 23:12:58 INFO db.DBInputFormat: Using read committed transaction isolation
18/07/04 23:12:59 INFO mapreduce.JobSubmitter: number of splits:1
18/07/04 23:12:59 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1530722670909_0004
18/07/04 23:13:00 INFO impl.YarnClientImpl: Submitted application
application_1530722670909_0004
18/07/04 23:13:00 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1530722670909_0004/
18/07/04 23:13:00 INFO mapreduce.Job: Running job: job_1530722670909_0004
18/07/04 23:13:12 INFO mapreduce.Job: Job job_1530722670909_0004 running in uber mode : false
18/07/04 23:13:12 INFO mapreduce.Job: map 0% reduce 0%
18/07/04 23:13:21 INFO mapreduce.Job: map 100% reduce 0%
18/07/04 23:13:22 INFO mapreduce.Job: Job job_1530722670909_0004 completed successfully
18/07/04 23:13:23 INFO mapreduce.Job: Counters: 30

File System Counters

FILE: Number of bytes read=0
FILE: Number of bytes written=127884
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=87
HDFS: Number of bytes written=37
HDFS: Number of read operations=4
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

Job Counters

Launched map tasks=1
Other local map tasks=1
Total time spent by all maps in occupied slots (ms)=6412
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=6412
Total vcore-milliseconds taken by all map tasks=6412
Total megabyte-milliseconds taken by all map tasks=6565888

Map-Reduce Framework

Map input records=1
Map output records=1
Input split bytes=87
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=147
CPU time spent (ms)=1620
Physical memory (bytes) snapshot=75149312
Virtual memory (bytes) snapshot=347095040
Total committed heap usage (bytes)=21573632

File Input Format Counters

Bytes Read=0

File Output Format Counters

Bytes Written=37

18/07/04 23:13:23 INFO mapreduce.ImportJobBase: Transferred 37 bytes in 28.5161 seconds (1.2975 bytes/sec)
18/07/04 23:13:23 INFO mapreduce.ImportJobBase: Retrieved 1 records.
[acadgild@localhost ~]\$ hdfs dfs -ls /SQOOPOUT1/
18/07/04 23:13:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-07-04 23:13 /SQOOPOUT1/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 37 2018-07-04 23:13 /SQOOPOUT1/part-m-00000
[acadgild@localhost ~]\$ hdfs dfs -cat /SQOOPOUT1/part-m-00000
18/07/04 23:14:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
3,James,Brown,New York,United States

[acadgild@localhost ~]\$ sqoop export --connect jdbc:mysql://localhost:3306/sqoop_export1 --table PersonExportNew --export-dir /SQOOPOUT1/part-m-00000 --username root -P
Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0../hcatalog does not exist!
HCatalog jobs will fail.
Please set \$HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.7.bin__hadoop-2.6.0../accumulo does not exist!
Accumulo imports will fail.
Please set \$ACCUMULO_HOME to the root of your Accumulo installation.
18/07/04 23:47:25 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
Enter password:
18/07/04 23:47:31 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/07/04 23:47:31 INFO tool.CodeGenTool: Beginning code generation
Wed Jul 04 22:47:31 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/07/04 23:47:32 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `PersonExport` AS t LIMIT 1
18/07/04 23:47:32 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `PersonExport` AS t LIMIT 1
18/07/04 23:47:32 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/9b67cb31fd26ee856e85ec798f15db9c/PersonExportNew.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/07/04 23:47:35 INFO orm.CompilationManager: Writing jar file:
/tmp/sqoop-acadgild/compile/9b67cb31fd26ee856e85ec798f15db9c/PersonExport.jar
18/07/04 22:47:35 INFO mapreduce.ExportJobBase: Beginning export of PersonExport
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in

[jar:file:/home/acadgild/install/hbase/hbase-1.4.4/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

18/07/04 23:47:36 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/07/04 23:47:36 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

18/07/04 23:47:38 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative

18/07/04 23:47:38 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

18/07/04 23:47:38 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

18/07/04 23:47:38 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

18/07/04 23:47:41 INFO input.FileInputFormat: Total input paths to process : 1

18/07/04 23:47:41 INFO input.FileInputFormat: Total input paths to process : 1

18/07/04 23:47:41 INFO mapreduce.JobSubmitter: number of splits:4

18/07/04 23:47:41 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

18/07/04 23:47:42 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_153072670909_0003

18/07/04 23:47:42 INFO impl.YarnClientImpl: Submitted application
application_1530722670909_0003

18/07/04 23:47:42 INFO mapreduce.Job: The url to track the job:

http://localhost:8088/proxy/application_1530722670909_0003/

18/07/04 23:47:42 INFO mapreduce.Job: Running job: job_1530722670909_0003

18/07/04 23:47:55 INFO mapreduce.Job: Job job_1530722670909_0003 running in uber mode : false

18/07/04 23:47:55 INFO mapreduce.Job: map 0% reduce 0%

18/07/04 23:48:24 INFO mapreduce.Job: map 50% reduce 0%

18/07/04 23:48:25 INFO mapreduce.Job: map 100% reduce 0%

18/07/04 23:48:25 INFO mapreduce.Job: Job job_1530722670909_0003 completed successfully

18/07/04 23:48:25 INFO mapreduce.Job: Counters: 30

File System Counters

FILE: Number of bytes read=0

FILE: Number of bytes written=511648

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=670

HDFS: Number of bytes written=0

HDFS: Number of read operations=19

HDFS: Number of large read operations=0

HDFS: Number of write operations=0

Job Counters

Launched map tasks=4

Data-local map tasks=4

Total time spent by all maps in occupied slots (ms)=110284

Total time spent by all reduces in occupied slots (ms)=0

Total time spent by all map tasks (ms)=110284
Total vcore-milliseconds taken by all map tasks=110284
Total megabyte-milliseconds taken by all map tasks=112930816
Map-Reduce Framework
Map input records=1
Map output records=1
Input split bytes=556
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=1524
CPU time spent (ms)=6320
Physical memory (bytes) snapshot=295948288
Virtual memory (bytes) snapshot=1389920256
Total committed heap usage (bytes)=68976640

File Input Format Counters

Bytes Read=0

File Output Format Counters

Bytes Written=0

18/07/04 23:48:25 INFO mapreduce.ExportJobBase: Transferred 670 bytes in 47.931 seconds (13.9784 bytes/sec)

18/07/04 23:48:25 INFO mapreduce.ExportJobBase: Exported 1 records.

[acadgild@localhost ~]\$ mysql -u root -p

Enter password:

Welcome to the MySQL monitor. Commands end with ; or \g.

Your MySQL connection id is 28

Server version: 8.0.11 MySQL Community Server - GPL

Copyright (c) 2000, 2018, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> use sqoop_export1;

Reading table information for completion of table and column names

You can turn off this feature to get a quicker startup with -A

Database changed

mysql> select * from PersonExportNew;

```
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area   | city      |
+-----+-----+-----+-----+-----+
|      3 | James | Brown | New York | United States |
+-----+-----+-----+-----+-----+
```

1 row in set (0.00 sec)

