

**Session 24:**

# **KAFKA INTRODUCTION**

**Assignment 2**

## Assignment 24: Apache Kafka II Assignment Problems

### Problem Statement

#### Data Sets Present:

#### 1. dataset\_producer.txt :

#### Terminal Execution :

```
[acadgild@localhost ~]$ cat /home/acadgild/Desktop/dataset_producer.txt
ItemTopic-{"item_id":"101"}-{"user_id":"U101"}
UserTopic-{"name":"John"}-{"exp":16}
ItemTopic-{"item_id":"101"}-{"user_id":"U106"}
UserTopic-{"name":"Mark"}-{"exp":18}
ItemTopic-{"item_id":"102"}-{"user_id":"U110"}
UserTopic-{"name":"Cylin"}-{"exp":15}
ItemTopic-{"item_id":"102"}-{"user_id":"U101"}
UserTopic-{"name":"Prod"}-{"exp":14}
ItemTopic-{"item_id":"104"}-{"user_id":"U102"}
UserTopic-{"name":"Abhay"}-{"exp":17}
ItemTopic-{"item_id":"107"}-{"user_id":"U104"}
UserTopic-{"name":"Misano"}-{"exp":19}[acadgild@localhost ~]$
```

### Task 1

Create a java program MyKafkaProducer.java that takes a file name and delimiter as input arguments.

It should read the content of file line by line.

Fields in the file are in following order

1. Kafka Topic Name
2. Key
3. value

For every line Created, insert the key and value to the repsective Kafka broker in a fire and forget mode.

**After record is sent, it should print appropriate message on screen.  
Pass dataset\_producer.txt as the input file and -as delimiter.**

### **Solution :**

#### **1. MyKafkaProducer.java :**

```
package com.acadgild.kafka.api;

import java.io.File;
import java.io.FileNotFoundException;
import java.util.ArrayList;
import java.util.List;
import java.util.Properties;
import java.util.Scanner;

import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;

public class MyKafkaProducer {

    public static void main(String[] args){

        List<String> lines = new ArrayList<String>();

        try {
            Scanner scanner = new Scanner(new
File("/home/acadgild/Desktop/dataset_producer.txt"));
            while (scanner.hasNextLine()) {
                //      System.out.println(scanner.nextLine());
                lines.add(scanner.nextLine());
            }
            scanner.close();
        } catch (FileNotFoundException e) {
            e.printStackTrace();
        }
    }
}
```

```

        Properties props = new Properties();
        props.put("bootstrap.servers", "localhost:9092");
        props.put("acks", "0");
        props.put("retries", 0);
        props.put("key.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
        props.put("value.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
        props.put("partitioner.class", "com.acadgild.kafka.api.RangePartitioner");

        KafkaProducer<String, String> producer = new KafkaProducer<>(props);
        ProducerRecord<String, String> producerRecord = null;

        for (String line : lines) {
            String[] token = line.split("-");
            producerRecord = new ProducerRecord<String, String>(token[0],
token[1],token[2]);
            producer.send(producerRecord);
            System.out.println(token[0] + " " +token[1]+" "+token[2]+" Sent!");
        }
        producer.close();
    }
}

```

## 2. RangePartitioner.java

```

package com.acadgild.kafka.api;

import java.util.List;
import java.util.Map;
import org.apache.kafka.clients.producer.Partitioner;
import org.apache.kafka.common.Cluster;
import org.apache.kafka.common.PartitionInfo;

public class RangePartitioner implements Partitioner{

    public void configure(Map<String, ?> configs) {}

    public int partition(String topic, Object key, byte[] keyBytes,
        Object value, byte[] valueBytes, Cluster cluster) {
        List<PartitionInfo> partitions =

```

```

cluster.partitionsForTopic(topic);

        // Getting the number of partitions for the topic
        int numPartitions = partitions.size();
        System.out.println("Topic: " + topic + "The key: " +
keyBytes.toString() + " value: " + valueBytes.toString() + "\n");
        return numPartitions - 1;

    }

    public void close() {}

}

```

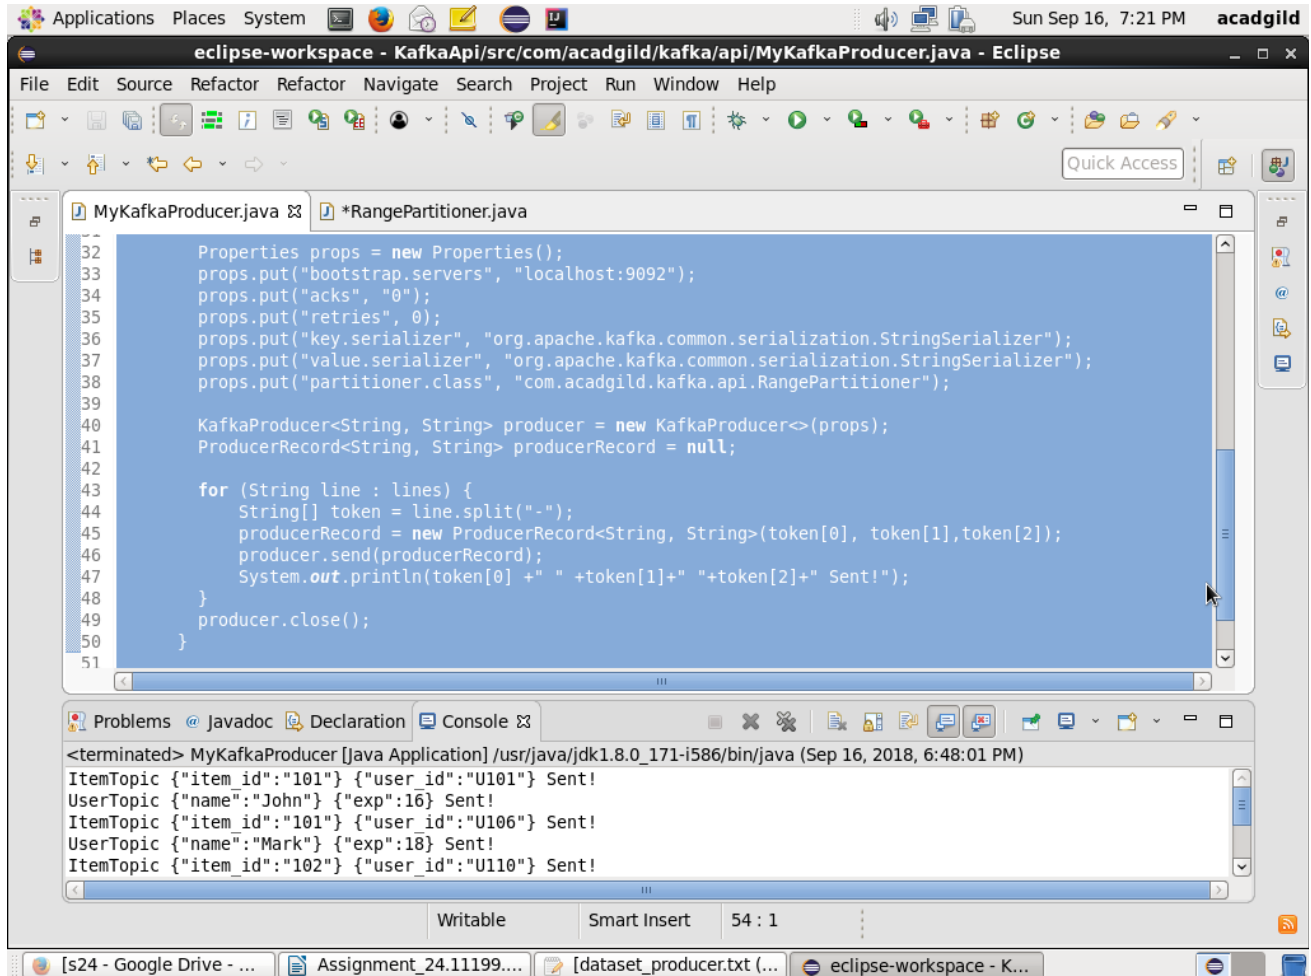
### **3. Result :**

```

ItemTopic {"item_id":"101"} {"user_id":"U101"} Sent!
UserTopic {"name":"John"} {"exp":16} Sent!
ItemTopic {"item_id":"101"} {"user_id":"U106"} Sent!
UserTopic {"name":"Mark"} {"exp":18} Sent!
ItemTopic {"item_id":"102"} {"user_id":"U110"} Sent!
UserTopic {"name":"Cylin"} {"exp":15} Sent!
ItemTopic {"item_id":"102"} {"user_id":"U101"} Sent!
UserTopic {"name":"Prod"} {"exp":14} Sent!
ItemTopic {"item_id":"104"} {"user_id":"U102"} Sent!
UserTopic {"name":"Abhay"} {"exp":17} Sent!
ItemTopic {"item_id":"107"} {"user_id":"U104"} Sent!
UserTopic {"name":"Misano"} {"exp":19} Sent!

```

## 4. Output :



The screenshot displays the Eclipse IDE interface. The main editor window shows the source code for `MyKafkaProducer.java`. The code configures a Kafka producer with properties like `bootstrap.servers`, `acks`, `retries`, `key.serializer`, `value.serializer`, and `partitioner.class`. It then iterates over a list of lines, splitting each line into tokens and sending them as `ProducerRecord` objects to a Kafka topic. The console output shows the execution results, indicating that several messages were successfully sent.

```
32 Properties props = new Properties();
33 props.put("bootstrap.servers", "localhost:9092");
34 props.put("acks", "0");
35 props.put("retries", 0);
36 props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
37 props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
38 props.put("partitioner.class", "com.acadgild.kafka.api.RangePartitioner");
39
40 KafkaProducer<String, String> producer = new KafkaProducer<>(props);
41 ProducerRecord<String, String> producerRecord = null;
42
43 for (String line : lines) {
44     String[] token = line.split("-");
45     producerRecord = new ProducerRecord<String, String>(token[0], token[1], token[2]);
46     producer.send(producerRecord);
47     System.out.println(token[0] + " " + token[1] + " " + token[2] + " Sent!");
48 }
49 producer.close();
50
51
```

Console Output:

```
<terminated> MyKafkaProducer [Java Application] /usr/java/jdk1.8.0_171-i586/bin/java (Sep 16, 2018, 6:48:01 PM)
ItemTopic {"item_id":"101"} {"user_id":"U101"} Sent!
UserTopic {"name":"John"} {"exp":16} Sent!
ItemTopic {"item_id":"101"} {"user_id":"U106"} Sent!
UserTopic {"name":"Mark"} {"exp":18} Sent!
ItemTopic {"item_id":"102"} {"user_id":"U110"} Sent!
```

## Task 2

**Modify the previous program MyKafkaProducer.java and create a new Java program KafkaProducerWithAck.java.**

**This should perform the same task as of KafkaProducer.java with some modification. When passing any data to a topic, it should wait for acknowledgement.**

**After acknowledgement is received from the broker, it should print the key and value which has been written to a specified topic.**

**The application should attempt for 3 retries before giving any exception. Pass dataset\_producer.txt as the input file and - as delimiter.**

## Solution :

### 1. MyKafkaProducerWithAck.java :

```
package com.acadgild.kafka.api;

import java.io.File;
import java.io.FileNotFoundException;
import java.util.ArrayList;
import java.util.List;
import java.util.Properties;
import java.util.Scanner;
import java.util.concurrent.ExecutionException;

import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.ProducerRecord;

public class MyKafkaProducerWithAck {

    public static void main(String[] args){

        List<String> lines = new ArrayList<String>();

        try {
            Scanner scanner = new Scanner(new
            File("/home/acadgild/Desktop/dataset_producer.txt"));
            while (scanner.hasNextLine()) {
                //      System.out.println(scanner.nextLine());
            }
        }
    }
}
```

```

        lines.add(scanner.nextLine());
    }
    scanner.close();
} catch (FileNotFoundException e) {
    e.printStackTrace();
}

```

```

Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("acks", "1");
props.put("retries", 3);
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
props.put("partitioner.class", "com.acadgild.kafka.api.RangePartitioner");

```

```

KafkaProducer<String, String> producer = new KafkaProducer<>(props);
ProducerRecord<String, String> producerRecord = null;

```

```

for (String line : lines) {
    String[] token = line.split("-");
    producerRecord = new ProducerRecord<String, String>(token[0], token[1], token[2]);
    int retries = 0;
    try {
        retries++;
        producer.send(producerRecord).get();
        System.out.println(token[0] + " " + token[1] + " " + token[2] + " Sent!");
    } catch (InterruptedException e) {
        // TODO Auto-generated catch block
        e.printStackTrace();
        if(retries >= 3)
            System.out.println("Error :" + e.getMessage());
    } catch (ExecutionException e) {
        // TODO Auto-generated catch block
        e.printStackTrace();
        if(retries >= 3)
            System.out.println("Error :" + e.getMessage());
    }
}

}
producer.close();
}

}

```



## 2. RangePartitioner.java

```
package com.acadgild.kafka.api;

import java.util.List;
import java.util.Map;
import org.apache.kafka.clients.producer.Partitioner;
import org.apache.kafka.common.Cluster;
import org.apache.kafka.common.PartitionInfo;

public class RangePartitioner implements Partitioner{

    public void configure(Map<String, ?> configs) {}

    public int partition(String topic, Object key, byte[] keyBytes,
        Object value, byte[] valueBytes, Cluster cluster) {
        List<PartitionInfo> partitions =
cluster.partitionsForTopic(topic);

        // Getting the number of partitions for the topic
        int numPartitions = partitions.size();
        System.out.println("Topic: " + topic + "The key: " +
keyBytes.toString() + " value: " + valueBytes.toString() + "\n");
        return numPartitions - 1;

    }

    public void close() {}

}
```

## 3. Result :

```
ItemTopic {"item_id":"101"} {"user_id":"U101"} Sent!
UserTopic {"name":"John"} {"exp":16} Sent!
ItemTopic {"item_id":"101"} {"user_id":"U106"} Sent!
UserTopic {"name":"Mark"} {"exp":18} Sent!
ItemTopic {"item_id":"102"} {"user_id":"U110"} Sent!
UserTopic {"name":"Cylin"} {"exp":15} Sent!
ItemTopic {"item_id":"102"} {"user_id":"U101"} Sent!
UserTopic {"name":"Prod"} {"exp":14} Sent!
ItemTopic {"item_id":"104"} {"user_id":"U102"} Sent!
UserTopic {"name":"Abhay"} {"exp":17} Sent!
ItemTopic {"item_id":"107"} {"user_id":"U104"} Sent!
UserTopic {"name":"Misano"} {"exp":19} Sent!
```